

# LAB ASSIGNMENT - 4

---

**Aim:** To study and learn about various system calls.

**To perform:** Comprehensive study of different categories of Linux system calls, categorized as.

## **1. Process Management System calls:**

**a) fork()** is a system call used to create a new process by duplicating the existing (parent) process. The new process created is called the child process.

When fork() is called, the entire memory space of the parent is duplicated. Both processes (parent and child) continue executing independently from the point where fork() was called.

It returns 0 in the child process and positive PID of the child in the parent process. It will return -1 on error (e.g., system can't create a new process).

**b) exec()** replaces the current process image with a new process image. It's used to run a new program from within a process.

There are multiple forms like:

- execl(), execlp()
- execv(), execvp(), execve()

Each differs in how arguments and the program path are passed:

- 'l' → arguments listed
- 'v' → arguments passed as vector (array)
- 'p' → search program in 'PATH'

It does not return any value if successful, it replaces the current process. If it fails, returns -1, and the old process continues.

**c) wait()** is used by a parent process to wait for its child process to terminate. It is used to prevent zombie processes (terminated child processes that haven't been cleaned up) and allows the parent to retrieve the exit status of the child.

It Returns the PID of the terminated child and stores exit status in the status variable.

## **2. File Management System calls**

### **a) open()**

The open() system call is used to open a file and obtain a file descriptor, which is a unique identifier for the opened file. This descriptor is then used in subsequent operations on the file<sup>24</sup>.

The syntax for open() is:

**int open(const char \*path, int flags, ... /\* mode\_t mode \*/);**

Key parameters include:

- path: The pathname of the file to be opened
- flags: Specify the access mode (O\_RDONLY, O\_WRONLY, O\_RDWR) and other options
- mode: Used only when creating a new file, specifies the permissions

The function returns the smallest available file descriptor, which is a non-negative integer. In case of an error, it returns -1

#### **b) read()**

The read() system call is used to read data from an open file into a memory buffer. It takes the file descriptor, buffer address, and the number of bytes to read as parameters<sup>12</sup>.

The typical signature of read() is:

```
ssize_t read(int fd, void *buf, size_t count);
```

Where:

- fd: The file descriptor obtained from open()
- buf: A pointer to the buffer where data will be stored
- count: The maximum number of bytes to read

The function returns the number of bytes actually read, which may be less than the requested count. A return value of 0 indicates end-of-file, and -1 indicates an error

#### **c) write()**

The write() system call writes data from a memory buffer to an open file. It takes the file descriptor, buffer address, and the number of bytes to write as parameters.

The typical signature is:

```
ssize_t write(int fd, const void *buf, size_t count);
```

Where:

- fd: The file descriptor obtained from open()
- buf: A pointer to the buffer containing data to be written
- count: The number of bytes to write

The function returns the number of bytes actually written, which may be less than the requested count. A return value of -1 indicates an error.

#### **d) close()**

The close() system call is used to close an open file descriptor, releasing the resources associated with it. Once a file descriptor is closed, it becomes available for reuse in subsequent open() calls<sup>1</sup>.

The typical signature is:

```
int close(int fd);
```

Where fd is the file descriptor to be closed. The function returns 0 on success and -1 on error.

Closing files is important because the operating system limits the number of files a process can have open simultaneously. Properly closing files prevents resource leaks and ensures that all data is properly written to disk.

### **3. Device Management System Calls**

Device management system calls are used to interact with hardware devices in Linux. These calls treat devices as files (following the "everything is a file" philosophy of Unix/Linux), but with special considerations for hardware interaction.

#### **read() and write() for Devices**

In Linux, device files represent physical devices and are located in the /dev directory. The same read() and write() system calls used for regular files are also used for device files, but their behavior depends on the specific device driver implementation<sup>5</sup>.

For character devices:

- read() retrieves data from the device into a buffer
- write() sends data from a buffer to the device

For example, when writing to a device file connected to a modem:

```
int fd = open("/dev/ttyS0", O_WRONLY);
```

```
write(fd, "AT\r", 3);
```

This approach allows applications to interact with hardware using the same familiar file I/O interface, maintaining consistency across the operating system

#### **c) ioctl()**

The ioctl() (input/output control) system call provides a way to send device-specific commands to device drivers. It is used when simple read/write operations are not sufficient for controlling a device's behavior.

The typical signature is:

```
int ioctl(int fd, unsigned long request, ...);
```

Where:

- fd: The file descriptor for the device
- request: A device-specific request code
- Additional arguments depending on the request

For example, terminal settings are often controlled with `ioctl()` calls:

```
struct termios term;

ioctl(fd, TCGETS, &term); // Get current terminal settings

term.c_lflag &= ~ECHO; // Disable echo

ioctl(fd, TCSETS, &term); // Apply new settings
```

The `ioctl()` call is essential for operations that don't fit into the simple read/write model, such as changing device modes, querying device status, or sending specialized commands<sup>5</sup>.

#### d) `select()`

The `select()` system call allows a program to monitor multiple file descriptors, waiting until one or more become ready for some type of I/O operation. This is particularly useful for implementing non-blocking I/O and managing multiple input/output streams concurrently.

The typical signature is:

```
c

int select(int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);
```

Where:

- `nfds`: The highest-numbered file descriptor plus 1
- `readfds`: Set of descriptors to check for reading
- `writefds`: Set of descriptors to check for writing
- `exceptfds`: Set of descriptors to check for exceptions
- `timeout`: Maximum time to wait

The `select()` call is commonly used in network servers, GUI applications, and any scenario where a program needs to handle multiple input sources without blocking on any single one.

## 4. Network Management System Calls

Network management system calls provide the foundation for network programming in Linux, enabling processes to communicate over networks using various protocols.

#### a) `socket()`

The `socket()` system call creates a communication endpoint (socket) for network communication. It is the first step in establishing network connections.

The typical signature is:

```
int socket(int domain, int type, int protocol);
```

Where:

- `domain`: Specifies the protocol family (AF\_INET for IPv4, AF\_INET6 for IPv6)

- **type:** Specifies the communication semantics (SOCK\_STREAM for TCP, SOCK\_DGRAM for UDP)
- **protocol:** Usually set to 0 for the default protocol

For example, creating a TCP/IP socket:

```
int sock_fd = socket(AF_INET, SOCK_STREAM, 0);
```

The function returns a file descriptor for the new socket, which is then used in subsequent network operations.

### **b) connect()**

The connect() system call establishes a connection between a socket and a specified address. It is used by client applications to connect to servers.

The typical signature is:

```
int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

Where:

- **sockfd:** The socket file descriptor
- **addr:** Points to a sockaddr structure containing the destination address
- **addrlen:** Size of the address structure

For example, connecting to a web server:

```
struct sockaddr_in server_addr;
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(80);
inet_pton(AF_INET, "192.168.1.1", &server_addr.sin_addr);
connect(sock_fd, (struct sockaddr *)&server_addr, sizeof(server_addr));
```

### **send() and recv()**

The send() and recv() system calls are used to transmit and receive data over connected sockets.

The send() function has the signature:

c

```
ssize_t send(int sockfd, const void *buf, size_t len, int flags);
```

Where:

- **sockfd:** The socket file descriptor
- **buf:** Points to the data to be sent
- **len:** Length of the data in bytes
- **flags:** Special behavior flags (usually 0)

The `recv()` function has the signature:

```
ssize_t recv(int sockfd, void *buf, size_t len, int flags);
```

Where:

- `sockfd`: The socket file descriptor
- `buf`: Buffer to store received data
- `len`: Maximum length of the buffer
- `flags`: Special behavior flags (usually 0)

These functions return the number of bytes sent or received, or -1 on error.

## **5. System Information Management System Calls**

System information management system calls provide processes with information about the system and their execution environment.

### **a) getpid()**

The `getpid()` system call returns the process ID of the calling process. This unique identifier is assigned by the kernel when the process is created.

The typical signature is:

```
pid_t getpid(void);
```

Example usage:

```
pid_t my_pid = getpid();  
printf("My process ID is: %d\n", my_pid);
```

### **b) getuid()**

The `getuid()` system call returns the real user ID of the calling process. This ID identifies the user who owns the process.

The typical signature is:

```
uid_t getuid(void);
```

Example usage:

```
uid_t my_uid = getuid();  
printf("My user ID is: %d\n", my_uid);
```

There is also a related call, `geteuid()`, which returns the effective user ID, potentially different from the real user ID when a program is running with `setuid` permissions.

### **c) gethostname()**

The `gethostname()` system call retrieves the hostname of the current machine.

The typical signature is:

```
int gethostname(char *name, size_t len);
```

Where:

- name: Buffer to store the hostname
- len: Size of the buffer

Example usage:

```
char hostname[256];  
gethostname(hostname, sizeof(hostname));  
printf("Hostname: %s\n", hostname);
```

#### **d) sysinfo()**

The sysinfo() system call returns information about system statistics and resources, such as memory usage and load averages.

The typical signature is:

```
int sysinfo(struct sysinfo *info);
```

Where info points to a sysinfo structure that will be filled with system information.

Example usage:

```
struct sysinfo si;  
sysinfo(&si);  
printf("Total RAM: %lu bytes\n", si.totalram);  
printf("Free RAM: %lu bytes\n", si.freeram);  
printf("System uptime: %ld seconds\n", si.uptime);  
.
```